

# SUMMA: scalable universal matrix multiplication algorithm

R. A. VAN DE GEIJN<sup>1</sup> AND J. WATTS<sup>2</sup>

<sup>1</sup>*Department of Computer Sciences, The University of Texas at Austin, Taylor Hall 2.124, Austin, TX 78712-1188, USA*

<sup>2</sup>*Scalable Concurrent Programming Laboratory, California Institute of Technology, Pasadena, CA 91125, USA*

## SUMMARY

**In the paper we give a straightforward, highly efficient, scalable implementation of common matrix multiplication operations. The algorithms are much simpler than previously published methods, yield better performance, and require less work space. MPI implementations are given, as are performance results on the Intel Paragon system. ©1997 by John Wiley & Sons, Ltd.**

## 1. INTRODUCTION

It seems somewhat strange to be writing a paper on parallel matrix multiplication almost two decades after commercial parallel systems first became available. One would think that by now we would be able to manage such an apparently straightforward task with simple, highly efficient implementations. Nonetheless, we appear to have gained new insight into this problem.

Different approaches proposed for matrix–matrix multiplication include 1D-systolic[1], 2D-systolic[1], Cannon’s algorithm[1,2], broadcast-multiply-roll[3,4], and the transpose algorithm[5]. Two recent efforts extend the work by Fox *et al.* to general meshes of nodes: the paper by Choi *et al.*[6] uses a two-dimensional block-wrapped (block-cyclic) data decomposition, while the papers by Huss-Lederman *et al.*[7,8] use a ‘virtual’ 2-D torus wrap data layout. Both these efforts report very good performance attained on the Intel Touchstone Delta, achieving a sizeable percentage of peak performance.

The method presented in our paper has the benefit of being more general, simpler and more efficient. We explain our algorithms for the case where the matrices to be multiplied, as well as the result, are block-mapped identically to nodes. However, we show how this restriction can be easily relaxed to achieve the wrapped decompositions mentioned above, as well as more general decompositions.

This paper makes a number of contributions: We present a simple approach and its scalability analysis. In addition, we give *complete* Message-Passing Interface (MPI)[9] implementations, demonstrating the power of this standard for coding concurrent algorithms. We show how our simpler approach outperforms more complex implementations, and, finally, we show how it is more general than alternative approaches to the problem.

After the original submission of this paper, we were made aware of the paper by Agarwal *et al.*[10], which presents essentially the algorithm we propose for  $C = AB$ . In their paper, they also carefully demonstrate the potential for overlapping of computation and communication, thereby providing a further optimization. In their paper, they do not address the formation of  $C = AB^T$  and the other variants, nor do they make the observation that

this approach is inherently appropriate in applications of the matrix–matrix multiplication like the LU factorization. Our own inspiration for the presented methods came from an *out-of-core* LU factorization implemented by Marques and van de Geijn in 1993. While this work was never published, in a subsequent out-of-core solver, we also used repeated rank- $k$  updates to implement the matrix–matrix multiply[11].

## 2. NOTATION

We consider the formation of the matrix products:

$$C = \alpha AB + \beta C \quad (1)$$

$$C = \alpha AB^T + \beta C \quad (2)$$

$$C = \alpha A^T B + \beta C \quad (3)$$

$$C = \alpha A^T B^T + \beta C \quad (4)$$

These are the special cases implemented as part of the widely used sequential Basic Linear Algebra Subprograms (BLAS)[12].

We will assume that each matrix  $X$  is of dimension  $m^X \times n^X$ ,  $X \in \{A, B, C\}$ . Naturally, there are constraints on these dimensions for the multiplications to be well-defined. We assume that the dimensions of  $C$  are  $m \times n$ , while the ‘other’ dimension is  $k$ .

## 3. MODEL OF COMPUTATION

We assume that the nodes of the parallel computer form an  $r \times c$  mesh. While for the analysis this is a physical mesh, the developed codes require only that the nodes can be logically configured as an  $r \times c$  mesh. The  $p = rc$  nodes are indexed by their row and column index, and the  $(i, j)$  node will be denoted by  $\mathbf{P}_{ij}$ .

In the absence of network conflicts, communicating a message between two nodes requires time  $\alpha + n\beta$ , which is reasonable on machines like the Intel Paragon system[13]. Parameters  $\alpha$  and  $\beta$  represent the startup cost and per item transfer time, respectively. Performing a floating point computation requires time  $\gamma$ .

## 4. DATA DECOMPOSITION

We will consider two dimensional data decompositions. The analysis will automatically include the one dimensional cases by letting either the row or column dimension of the mesh equal one. For all the algorithms we assume the following assignment of data to nodes. Given  $m^X \times n^X$  matrix  $X$ ,  $X \in \{A, B, C\}$ , and an  $r \times c$  logical mesh of nodes, we partition as follows:

$$X = \left( \begin{array}{c|c|c} X_{00} & \cdots & X_{0(c-1)} \\ \vdots & & \vdots \\ \hline X_{(r-1)0} & \cdots & X_{(r-1)(c-1)} \end{array} \right)$$

and assign  $X_{ij}$  to node  $\mathbf{P}_{ij}$ . Submatrix  $X_{ij}$  has dimensions  $m_i^X \times n_j^X$ , with  $\sum m_i^X = m^X$  and  $\sum n_j^X = n^X$ . Further restrictions on the dimensions are given for each of the four variants, to ensure that the appropriate row and column dimensions match.

## 5. FORMING $C = \alpha AB + \beta C$

For this operation to be well-defined, we require  $m^A = m$ ,  $n^A = m^B = k$ , and  $n^B = n$ . For simplicity, we take  $\alpha = 1$  and  $\beta = 0$  in our description. If  $a_{ij}$ ,  $b_{ij}$  and  $c_{ij}$  denote the  $(i, j)$  element of the matrices, respectively, then the elements of  $C$  are given by

$$c_{ij} = \sum_{l=0}^{k-1} a_{il} b_{lj}$$

Note that rows of  $C$  are computed from rows of  $A$ , and columns of  $C$  are computed from columns of  $B$ . We hence restrict our data decomposition so that rows of  $A$  and  $C$  are assigned to the same row of nodes, and columns of  $B$  and  $C$  are assigned to the same column of nodes. Hence,  $m_i^C = m_i^A$  and  $n_j^C = n_j^B$ .

### 5.1. Basic parallel algorithm

Let us consider what computation is required to form  $C_{ij}$ :

$$C_{ij} = \left( \overbrace{A_{i0} \mid A_{i1} \mid \cdots \mid A_{i(c-1)}}^{\tilde{A}_i} \right) \left( \begin{array}{c} B_{0j} \\ B_{1j} \\ \vdots \\ B_{(r-1)j} \end{array} \right) \tilde{B}^j$$

Notice that  $\tilde{A}_i$  is entirely assigned to node row  $i$ , while  $\tilde{B}^j$  is entirely assigned to node column  $j$ . Letting

$$\tilde{A}_i = \left( \tilde{a}_i^0 \mid \tilde{a}_i^1 \mid \cdots \mid \tilde{a}_i^{k-1} \right) \text{ and } \tilde{B}^j = \left( \begin{array}{c} \tilde{b}_0^{jT} \\ \tilde{b}_1^{jT} \\ \vdots \\ \tilde{b}_{k-1}^{jT} \end{array} \right)$$

we see that

$$C_{ij} = \sum_{l=0}^{k-1} \tilde{a}_i^l \tilde{b}_l^{jT}$$

Hence the matrix-matrix multiply can be formulated as a sequence of rank-one updates.

It now suffices to parallelize each rank-one update. Pseudo-code for this, executed simultaneously on all nodes  $\mathbf{P}_{ij}$  is given in Figure 1. The process is illustrated in Figure 2.


To analyze the cost of this basic algorithm, we make some simplifying assumptions:  $m_i^C = m_i^A = m/r$ ,  $n_i^C = n_i^B = n/r$ ,  $n_i^A = k/c$ , and  $m_i^B = k/r$ . Since relatively little data is involved during each broadcast, we will assume a minimum spanning tree broadcast is used and the cost of our algorithm is given by

$$k \left[ \frac{2mn}{p} \gamma + \lceil \log(c) \rceil \left( \alpha + \frac{m}{r} \beta \right) + \lceil \log(r) \rceil \left( \alpha + \frac{n}{c} \beta \right) \right]$$

```

 $C_{ij} = 0$ 
for  $l = 0, k - 1$ 
    broadcast  $\tilde{a}_i^l$  within my row
    broadcast  $\tilde{b}_l^j$  within my column
     $C_{ij} = C_{ij} + \tilde{a}_i^l \tilde{b}_l^j T$ 
endfor

```

Figure 1. Pseudo-code for  $C = AB$ Figure 2. Operations implementing the inner loop of Figure 1 of a  $2 \times 3$  mesh of nodes

The terms within the square brackets are due to the rank-one update, broadcast within row, and broadcast within column, respectively. We ignore the packing required before sending rows of  $\tilde{B}^i$ . The total time is thus

$$T(m, n, k, p) = \frac{2mnk}{p} \gamma + k([\log(c)] + [\log(r)])\alpha + [\log(c)] \frac{mk}{r} \beta + [\log(r)] \frac{nk}{c} \beta \quad (5)$$

This compares to a sequential time of  $2mnk\gamma$ .

To establish the scalability of this approach, we analyse the case where  $m = n = k$ ,  $r = c = \sqrt{p}$ , and  $p$  is a power of two. Given the complexity in (5), the estimated speedup is

$$S(n, p) = \frac{2n^3\gamma}{\frac{2n^3}{p}\gamma + n \log(p)\alpha + \log(p) \frac{n^2}{\sqrt{p}}\beta} = \frac{p}{1 + \frac{p \log(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p} \log(p)}{2n} \frac{\beta}{\gamma}}$$

The corresponding efficiency is

$$E(n, p) = \frac{S(n, p)}{p} = \frac{1}{1 + \frac{p \log(p)}{2n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p} \log(p)}{2n} \frac{\beta}{\gamma}} = \frac{1}{1 + O\left(\frac{p \log(p)}{n^2}\right) + O\left(\frac{\sqrt{p} \log(p)}{n}\right)}$$

Ignoring the  $\log(p)$  term, which grows *very* slowly when  $p$  is reasonably large, we notice the following. If we increase  $p$  and we wish to maintain efficiency, we must increase  $n$  with  $\sqrt{p}$ . Since memory requirements grow with  $n^2$ , and physical memory grows linearly with  $p$  as nodes are added, we conclude that the method is scalable in the following sense. If we maintain memory use *per node*, this algorithm will maintain efficiency, if  $\log(p)$  is treated as a constant.

Alternative broadcast algorithms, e.g. pipelined or scatter-collect broadcasts [14,15], can be used to eliminate the  $\log(p)$  factor, at the expense of a larger number of startups. We instead will present the benefits of pipelining *computation and communication*.

## 5.2. Pipelined algorithm

Let us consider implementing the broadcast as passing of a message around the logical ring that forms the row or column. In this case, the time complexity becomes:

$$(c-1) \left( \alpha + \frac{m}{r} \beta \right) + (r-1) \left( \alpha + \frac{n}{c} \beta \right) \quad (6)$$

$$+ k \left( \frac{2mn}{p} \gamma + \alpha + \frac{m}{r} \beta + \alpha + \frac{n}{c} \beta \right) \quad (7)$$

$$+ (c-2) \left( \alpha + \frac{m}{r} \beta \right) + (r-2) \left( \alpha + \frac{n}{c} \beta \right) \quad (8)$$

$$+ \frac{2mn}{p} \gamma \quad (9)$$

$$= \frac{2mn(k+1)}{p} \gamma + (k+2c-3) \left( \alpha + \frac{m}{r} \beta \right) + (k+2r-3) \left( \alpha + \frac{n}{c} \beta \right) \quad (10)$$

Contribution (6) equals the time required for both the first column of  $\tilde{A}_{r-1}$  and the first row of  $\tilde{B}^{c-1}$  to reach  $\mathbf{P}_{(r-1)(c-1)}$  (filling the pipe); (7) equals the time for performing the local update and passing the messages; (8) equals the time for the final messages (initiated at  $\mathbf{P}_{(r-1)(c-1)}$ ) to reach the end of the pipe; and (9) equals the time for the final update at the node at the end of the pipe ( $\mathbf{P}_{(r-1)(c-2)}$  or  $\mathbf{P}_{(r-2)(c-1)}$ ). Note that, for large  $k$ , the ‘log’ factors in (5) are essentially removed.

To establish the scalability of the pipelined approach, we again analyse the case where  $m = n = k$  and  $r = c = \sqrt{p}$ . This changes the complexity in (10) to approximately

$$\frac{2n^3}{p} \gamma + 2(n + 2\sqrt{p} - 3) \left( \alpha + \frac{n}{\sqrt{p}} \beta \right)$$

and the estimated speedup is

$$S(n, p) = \frac{2n^3 \gamma}{\frac{2n^3}{p} \gamma + 2(n + 2\sqrt{p} - 3) \left( \alpha + \frac{n}{\sqrt{p}} \beta \right)} \approx \frac{p}{1 + \frac{p}{n^2} \frac{\alpha}{\gamma} + \frac{\sqrt{p}}{n} \frac{\beta}{\gamma}}$$

The corresponding efficiency is

$$E(n, p) = \frac{S(n, p)}{p} = \frac{1}{1 + O\left(\frac{p}{n^2}\right) + O\left(\frac{\sqrt{p}}{n}\right)}$$

The  $\log(p)$  term has disappeared and the method is again scalable in the sense that if we maintain memory use *per node*, this algorithm will maintain efficiency.

### 5.3. Blocking

Further improvements can be obtained by observing that reformulating the method in terms of matrix–matrix multiplications instead of rank-one updates can greatly improve the performance of an individual node. Matrix–matrix operations perform  $O(n^3)$  computation on  $O(n^2)$  data, thereby overcoming the memory bandwidth bottleneck present on most modern microprocessors. Highly optimized versions of an important set of such operations (the level-3 BLAS[12]) are typically provided by major vendors of high performance microprocessors.

We can stage the computation using matrix–matrix multiplication by accumulating several columns of  $\tilde{A}_i$  and rows of  $\tilde{B}_j$  before updating the local matrix. In our explanation, each column  $\tilde{a}_i^l$  becomes a *panel* of columns, and row  $\tilde{b}_j^l$  a corresponding *panel* of rows.

An additional advantage of blocking is that it reduces the number of messages incurred, thereby reducing communication overhead.

### 5.4. Code

MPI code for the pipelined, blocked algorithm is given in Figures 3 and 4. In the algorithm, arrays `a`, `b` and `c` hold the local matrices in column-major order. Parameters `lda`, `ldb` and `ldc` indicate the local leading dimension for arrays `a`, `b` and `c`, respectively. Variables `m`, `n`, and `k` hold  $m$ ,  $n$ , and  $k$ , respectively. Entries `m_a[i]`, `n_a[i]`, `m_b[i]`, `n_b[i]`, `m_c[i]` and `n_c[i]` hold  $m_i^A$ ,  $n_i^A$ ,  $m_i^B$ ,  $n_i^B$ ,  $m_i^C$  and  $n_i^C$ , respectively. Variable `nb` indicates the number of columns of  $\tilde{A}_i$  and rows of  $\tilde{B}^j$  that are accumulated before updated the local block of  $C$ . There are two work arrays: `work1` and `work2`, in which `nb` columns of  $\tilde{A}_i$  and `nb` rows of  $\tilde{B}^j$ , respectively, can be accumulated. MPI communicators indicating the nodes that constitute a row and column in the logical node mesh are provided in `comm_row` and `comm_col`.

Whenever possible, BLAS calls are used, and we use the LAPACK utility routine `d1acpy` to copy matrices to matrices.

## 6. FORMING $C = \alpha A^T B^T + \beta C$

Forming  $C = \alpha A^T B^T + \beta C$  can be easily derived by noting that  $C^T = BA$  and reversing the roles of rows and columns in the algorithm given in Section 5. Note that this computes  $C^T$ , which would need to be transposed. This transpose operation is a topic in itself. In a paper[16] written since we originally submitted this paper, we give insight into how this transpose operation can be avoided.

---

```

#include "mpi.h"
/* macro for column major indexing */
#define A( i,j ) (a[ j*lda + i ])
#define B( i,j ) (b[ j*ldb + i ])
#define C( i,j ) (c[ j*ldc + i ])

#define min( x, y ) ( (x) < (y) ? (x) : (y) )

int i_one=1; /* used for constant passed to blas call */
double d_one=1.0,
       d_zero=0.0; /* used for constant passed to blas call */

void pdgemm( m, n, k, nb, alpha, a, lda, b, ldb,
            beta, c, ldc, m_a, n_a, m_b, n_b, m_c, n_c,
            comm_row, comm_col, work1, work2 )
int m, n, k, /* global matrix dimensions */
    nb, /* panel width */
    m_a[], n_a[], /* dimensions of blocks of A */
    m_b[], n_b[], /* dimensions of blocks of A */
    m_c[], n_c[], /* dimensions of blocks of A */
    lda, ldb, ldc; /* leading dimension of local arrays that
                  hold local portions of matrices A, B, C */
double *a, *b, *c, /* arrays that hold local parts of A, B, C */
       alpha, beta, /* multiplication constants */
       *work1, *work2; /* work arrays */
MPI_Comm comm_row, /* communicator for this row of nodes */
         comm_col; /* communicator for this column of nodes */
{
    int myrow, mycol, /* my row and column index */
        nrow, ncol, /* number of node rows and columns */
        i, j, kk, iwrk, /* misc. index variables */
        icurrow, icurcol, /* index of row and column that hold current
                          row and column, resp., for rank-1 update */
        ii, jj; /* local index (on icurrow and icurcol, resp.)
                of row and column for rank-1 update */
    double *temp; /* temporary pointer used in pdgemm_abt */
    double *p;

    /* get myrow, mycol */
    MPI_Comm_rank( comm_row, &myrow ); MPI_Comm_rank( comm_col, &mycol );
    /* scale local block of C */
    for ( j=0; j<n_c[ mycol ]; j++ )
        for ( i=0; i<m_c[ myrow ]; i++ )
            C( i,j ) = beta * C( i,j );
}

```

Figure 3. MPI code for  $C = \alpha AB + \beta C$  and  $C = \alpha AB^T + \beta C$

## 7. FORMING $C = \alpha AB^T + \beta C$

One approach to implementing this algorithm is to transpose matrix  $B$  followed by the algorithm presented in Section 6. We will show how to avoid this initial communication.

For this operation to be well-defined, we require  $n^A = n$ ,  $m^A = m^B = k$  and  $n^B = m$ . Again, we take  $\alpha = 1$  and  $\beta = 0$  in our description. If  $a_{ij}$ ,  $b_{ij}$  and  $c_{ij}$  denote the  $(i, j)$  element of the matrices, respectively, then the elements of  $C$  are given by

$$c_{ij} = \sum_{l=0}^{k-1} a_{il} b_{jl}$$

```

icurrow = 0;          icurcol = 0;
ii = jj = 0;
/* malloc temp space for summation
temp = (double *) malloc(m_c[myrow]*nb*sizeof(double) );
NN USATÓ QUI

for ( kk=0; kk<k; kk+=iwrk) {
iwrk = min( nb, m_b[ icurrow ]-ii );
iwrk = min( iwrk, n_a[ icurcol ]-jj );
/* pack current iwrk columns of A into work1
if ( mycol == icurcol )
dlacpy_( "General", &m_a[ myrow ], &iwrk, &A( 0, jj ), &lda, work1,
&m_a[ myrow ] );
/* pack current iwrk rows of B into work2
if ( myrow == icurrow )
dlacpy_( "General", &iwrk, &n_b[ mycol ], &B( ii, 0 ), &ldb, work2,
&iwrk );
/* broadcast work1 and work2
RING_Bcast( work1 , m_a[ myrow ]*iwrk, MPI_DOUBLE, icurcol, comm_row );
RING_Bcast( work2 , n_b[ mycol ]*iwrk, MPI_DOUBLE, icurrow, comm_col );
/* update local block
dgemm_( "No transpose", "No transpose", &m_c[ myrow ], &n_c[ mycol ],
&iwrk, &alpha, work1, &m_b[ myrow ], work2, &iwrk, &d_one,
c, &ldc );
/* update icurcol, icurrow, ii, jj
ii += iwrk; jj += iwrk;
if ( jj>=n_a[ icurcol ] ) { icurcol++; jj = 0; };
if ( ii>=m_b[ icurrow ] ) { icurrow++; ii = 0; };
}
free( temp );
}

RING_Bcast( double *buf, int count, MPI_Datatype type, int root,
MPI_Comm comm )
{
int me, np;
MPI_Status status;

MPI_Comm_rank( comm, me ); MPI_Comm_size( comm, np );
if ( me != root )
MPI_Recv( buf, count, type, (me-1+np)%np, MPI_ANY_TAG, comm );
if ( ( me+1 )%np != root )
MPI_Send( buf, count, type, (me+1)%np, 0, comm );
}

```

Figure 4. MPI code for  $C = \alpha AB + \beta C$  continued; see Figure 3 for first part of code

For our algorithm, we make the restriction that columns of  $A$  and columns of  $B$  are assigned to the same column of nodes, and rows of  $A$  and rows of  $C$  are assigned to the same row of nodes.

### 7.1. Basic algorithm

Let

$$\tilde{C}_i = ( C_{i0} \mid C_{i1} \mid \cdots \mid C_{i(c-1)} ) = ( \tilde{c}_i^0 \mid \tilde{c}_i^1 \mid \cdots \mid \tilde{c}_i^{k-1} )$$



and

$$\tilde{B}^j = \begin{pmatrix} \frac{\tilde{b}_0^{jT}}{\tilde{b}_1^{jT}} \\ \vdots \\ \tilde{b}_{k-1}^{jT} \end{pmatrix}$$

Then matrix algebra tells us that

$$c_i^l = \sum_{j=0}^{c-1} A_{ij} \tilde{b}_l^{jT}$$

Since rows of  $C_{ij}$  and  $A_{ij}$  are assigned to the same row of nodes, and elements of  $\tilde{b}_l^j$  and columns of  $A_{ij}$  are assigned to the same column of nodes, we derive the pseudo-code given in Figure 5 for node  $P_{ij}$ . A picture describing the mechanism is given in Figure 6.

```

 $C_{ij} = 0$ 
for  $l = 0, k - 1$ 
    broadcast  $\tilde{b}_l^j$  within my column
    form  $\tilde{c}_i^{l,j} = A_{ij} \tilde{b}_l^{jT}$ 
    sum all  $\tilde{c}_i^{l,j}$  within my row to the
    node that holds  $\tilde{c}_i^l$ 
endfor

```

Figure 5. Pseudo-code for  $C = AB^T$

Assuming minimum spanning tree broadcast and sum-to-one are used, the cost becomes

$$k \left[ \lceil \log(c) \rceil \left( \alpha + \frac{m}{r} \beta \right) + \frac{2kn}{p} \gamma + \lceil \log(r) \rceil \left( \alpha + \frac{k}{c} \beta + \frac{n}{c} \gamma \right) \right] \quad (11)$$

Scalability properties are much like those of the algorithm for  $C = AB$ .

## 7.2. Pipelining

As with forming  $C = AB$ , the above algorithm can be improved by introducing pipelining. Let us consider implementing the broadcast as a passing of the message around the logical ring that forms the column. Similarly, let the summation within rows be implemented as a passing of a ‘bucket’ that collects all local contributions to the node on which the result is required. The effect on the time complexity is much like that obtained for the formation of  $C = AB$ .

## 7.3. Blocking

As for the computation of  $C = AB$ , further improvements can be obtained by observing that reformulating the method in terms of matrix–matrix multiplications instead of matrix–




Figure 6. Operations implementing the inner loop of Figure 5 of a  $2 \times 3$  mesh of nodes

vector multiplications can greatly improve the performance of an individual node. This can be accomplished by taking both  $\tilde{c}_i^l$  and  $\tilde{b}_l^j$  to be a small number of columns and rows, respectively. Again, communication overhead is reduced as well.

#### 7.4. Code

MPI code for the pipelined, blocked algorithm is given in Figures 3 and 7. The parameters are essentially the same as those used for forming  $C = \alpha AB + \beta C$ .

#### 8. FORMING $C = \alpha A^T B + \beta C$

Forming  $C = \alpha A^T B + \beta C$  can be easily derived by noting that  $C^T = \alpha B A^T + \beta C^T$  and reversing the roles of rows and columns in the algorithm given in Section 7.

---

```

icurrow = 0;          icurcol = 0;
ii = jj = 0;
/* malloc temp space for summation */
temp = (double *) malloc( m_c[myrow]*nb*sizeof(double) );
/* loop over all column panels of C */
for ( kk=0; kk<k; kk+=iwrk) {
    iwrk = min( nb, m_b[ icurrow ]-ii );
    iwrk = min( iwrk, n_c[ icurcol ]-jj );
    /* pack current iwrk rows of B into work2 */
    if ( myrow == icurrow )
        dlacpy_( "General", &iwrk, &n_b[ mycol ], &B( ii, 0 ), &ldb, work2,
                &iwrk );
    /* broadcast work2 */
    RING_Bcast( work2, n_b[ mycol ]*iwrk, MPI_DOUBLE, icurrow, comm_col );
    /* Multiply local block of A times incoming rows of B */
    dgemm_( "No transpose", "Transpose", &m_c[ myrow ], &iwrk,
            &n_a[ mycol ], &alpha, a, &lda, work2, &iwrk, &d_zero,
            work1, &m_c[ myrow ] );
    /* Sum to node that holds current columns of C */
    RING_SUM( work1, m_c[ myrow ]*iwrk, MPI_DOUBLE, icurcol, comm_row, temp );
    /* Add to current columns of C */
    if ( mycol == icurcol ) {
        p = work1;
        for ( j=jj; j<jj+iwrk; j++) {
            daxpy_( &m_c[ myrow ], &d_one, p, &i_one, &C( 0,j ), &i_one );
            p += m_c[ myrow ];
        }
    }
    /* update icurcol, icurrow, ii, jj */
    ii += iwrk;      jj += iwrk;
    if ( jj>=n_c[ icurcol ] ) { icurcol++; jj = 0; };
    if ( ii>=m_b[ icurrow ] ) { icurrow++; ii = 0; };
}
free( temp );
}

RING_SUM( double *buf, int count, MPI_Datatype type, int root,
          MPI_Comm comm, double *work )
{
    int me, np;
    MPI_Status status;

    MPI_Comm_rank( comm, &me );      MPI_Comm_size( comm, &np );
    if ( me != (root+1)%np ) {
        MPI_Recv( work, count, type, (me-1+np)%np, MPI_ANY_TAG, comm, &status );
        daxpy_( &count, &d_one, work, &i_one, buf, &i_one );
    }
    if ( me != root )
        MPI_Send( buf, count, type, (me+1)%np, 0, comm );
}

```

Figure 7. MPI code for  $C = \alpha AB^T + \beta C$  continued; see Figure 3 for first part of code

## 9. PERFORMANCE RESULTS

We do not compare the achieved performance with predicted performance. The reason for this is that there are too many parameters that cannot be easily controlled. For example, the performance of the BLAS kernel `dgemm`, which is used for the matrix–matrix multiplication on a single node, is highly dependent on the matrix size and other circumstances. Instead, we




Figure 8. Performance of SUMMA vs. PUMMA for  $C = AB$  on 64 nodes

compare the performance of our basic matrix-multiplication algorithms for  $C = \alpha AB + \beta C$  and  $C = \alpha AB^T + \beta C$  with that achieved by the PUMMA implementation in [6], which we obtained from `netlib`. We modified the PUMMA code to call the most efficient forms of NX communication primitives (including forced messages).

The implementation of SUMMA is essentially the one given in this paper. We implemented it using the MPI send and receive primitives used in the algorithms in this paper. However, the current MPI implementation on the Paragon incurs high latency and achieves lower bandwidth than the equivalent native NX calls. For that reason, we also implemented a highly optimized version that uses ‘forced’ (ready-receive) messages and NX calls.

In Figures 8 and 9, we report the performance *per node* for  $C = AB$  as a function of problem size for the two SUMMA versions and PUMMA. Peak performance observed for `dgemm` on a single node is in the 45 MFLOPS range. The blocking size (`nb`) chosen for both PUMMA and SUMMA was 20, which appears to give the best performance. Much more dramatic is the difference between PUMMA and SUMMA on a non-power-two mesh, results of which are reported in Figure 10. The primary reason is that the broadcast-multiply-roll algorithm generalizes more readily to non-square meshes when either the row or column dimension is an integer multiple of the other dimension. When this is not the case, performance suffers dramatically. SUMMA does not have the same kind of dependency, and it performs well.

In Figure 11 and 12, we report the performance *per node* as a function of the number of nodes when memory use per node is held constant for  $C = AB$ . Note that performance can essentially be maintained, as predicted by our analysis. Also, the cases we present represent

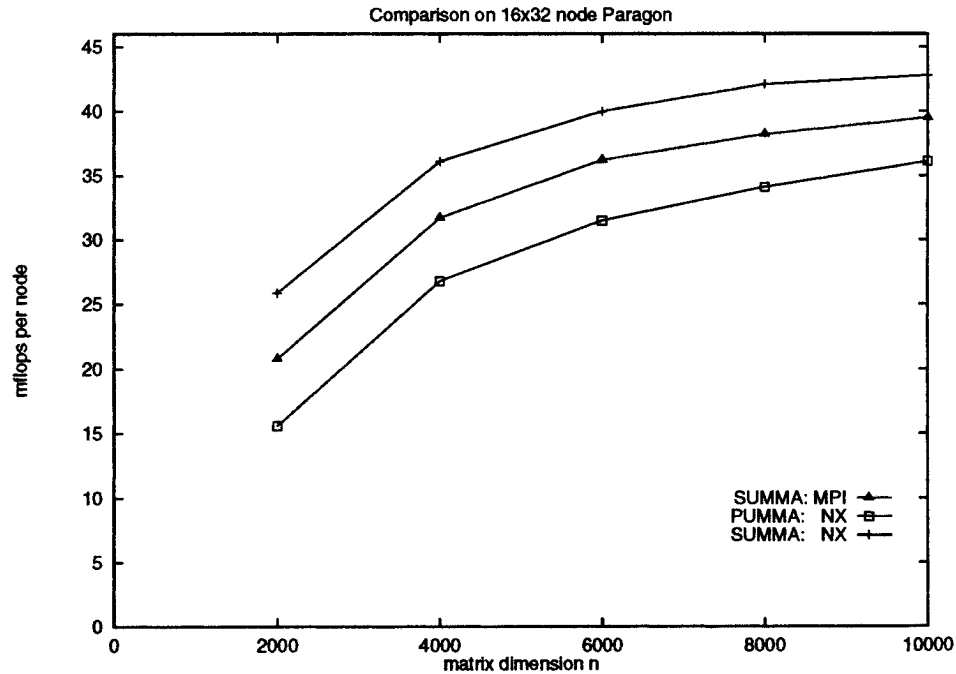


Figure 9. Performance of SUMMA vs. PUMMA for  $C = AB$  on 512 nodes

relatively small problems: a local  $500 \times 500$  problem requires only 2 Mbytes of memory. We need space for three of these matrices, plus a small amount of workspace. The Paragon typically comes with at least 32 Mbytes of memory per node.

In Figure 13, we report the performance per node for  $C = AB^T$  as a function of problem size for SUMMA and PUMMA. Peak performance on a single node is again in the 45 MFLOPS range. The observed performance is very similar to that of  $C = AB$ . Again, SUMMA is competitive with PUMMA. In Figure 14, we show the scalability of this algorithm.

It should be noted that PUMMA requires considerably more workspace per node than the amount of memory used to store the local matrices. SUMMA by comparison uses a lower order amount of memory per node, which means that SUMMA has the added advantage of allowing larger problems to be run.

## 10. GENERALIZATION OF THE CODE

The codes presented in this paper are slight simplifications, allowing us to include them in the body of the paper. Note that the following generalizations are easily obtained.

### 10.1. Relaxing the dependence on multiples of nb

We will limit our discussion in this Section to the forming of  $C = AB$  only. The same techniques can be easily applied to the other cases.

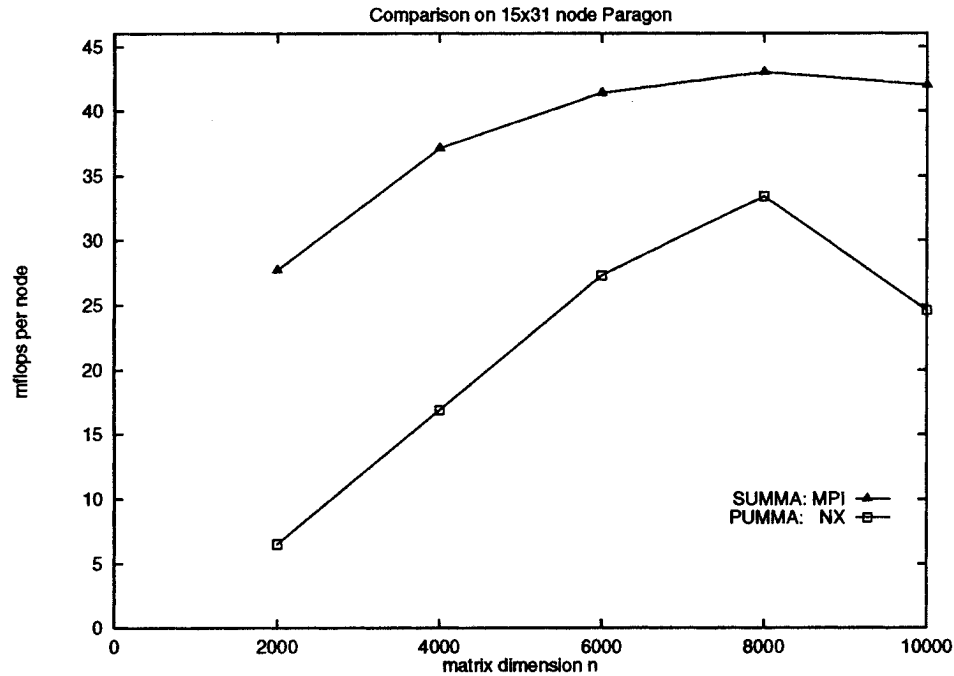


Figure 10. Performance of SUMMA vs. PUMMA for  $C = AB$  on 465 nodes

Note that the code requires nodes to have a multiple of  $nb$  rows and columns of the matrix, except for the last row and column of nodes. There are two alternatives for overcoming this restriction:

- The width of the current row and column panels can be reduced when a border between nodes is encountered.
- Partial panels can be passed to the next row or column of nodes, to be filled fully before broadcasting.

The latter is particularly easy to incorporate into the pipelined version of the algorithm.

## 10.2. Relaxing the alignment of the matrices

Note that it suffices that the same rows of  $A$  and  $C$  are assigned to the same row of nodes, and the same columns of  $B$  and  $C$  to the same column of nodes. If  $A$  and  $B$  are not aligned with  $C$  in the other dimension, it is a matter of passing in `icurrow` and `icurcol` as parameters to the routine.

## 10.3. Generalizing the matrix decomposition

We will further illustrate the flexibility of our approach by showing how the algorithm in Figure 4 can be easily changed to handle block-wrapped data decompositions like those used for ScaLAPACK. Indeed, it is a matter of changing the code segments

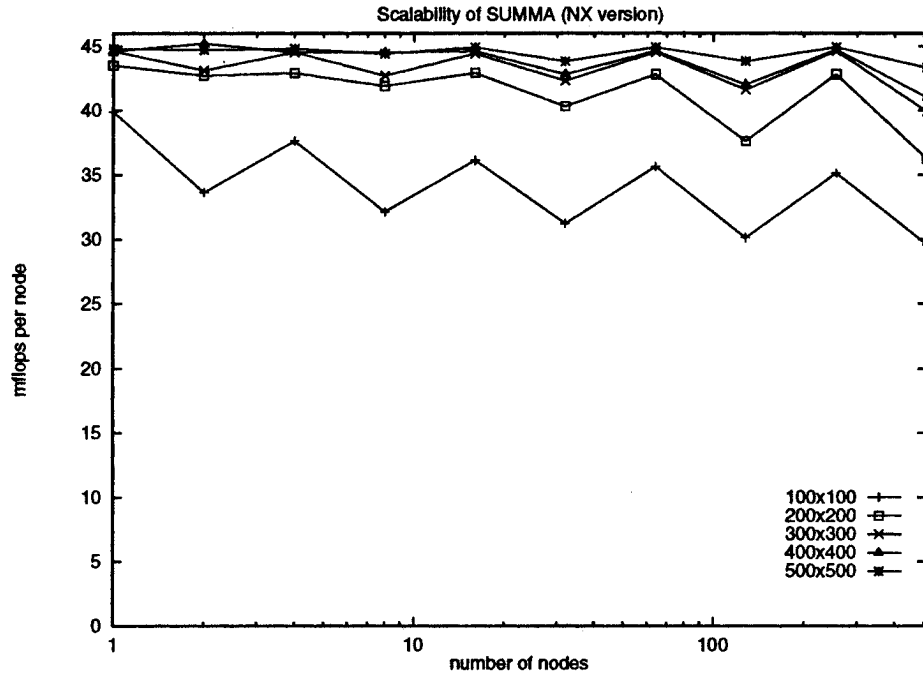


Figure 11. Performance of SUMMA (NX version) for  $C = AB$  as a function of the number of nodes, when memory use per node is held constant. The curves for  $100 \times 100$ ,  $200 \times 200$ , etc., indicate the use of local memory equivalent to that necessary to store a matrix of the indicated size, e.g., on 64 and 128 nodes,  $100 \times 100$  is equivalent to a  $800 \times 800$  and  $1128 \times 1128$  global matrix, respectively. The 'zigzagging' of the curves is due to the effects of square vs. non-square meshes

```

iwrk = min( nb, m_b[ icurrow ]-ii );
iwrk = min( iwrk, n_a[ icurcol ]-jj );

to

iwrk = min( nb, k-kk );

and

/* update icurcol, icurrow, ii, jj */
ii += iwrk;          jj += iwrk;
if ( jj > n_a[ icurcol ] ) { icurcol++; jj = 0; };
if ( ii > m_b[ icurrow ] ) { icurrow++; ii = 0; };

to

/* update icurcol, icurrow, ii, jj */
if ( myrow == icurrow ) ii += iwrk;
if ( mycol == icurcol ) jj += iwrk;
icurrow = ( icurrow+1 )%nprow;
icurcol = ( icurcol+1 )%npcol;

```

Other alternative matrix decompositions can be handled similarly.




Figure 12. Performance of PUMMA for  $C = AB$  as a function of the number of nodes, when memory use per node is held constant

#### 10.4. Odd-shaped matrices

A frequent use of matrix–matrix multiplication in applications is the case where  $k$  is much smaller than  $m$  and  $n$ . Examples of this occur in ScaLAPACK routines like those for the  $LU$  and  $QR$  factorization[17]. In such cases, our approach continues to be useful. However it may be necessary to substitute a minimum spanning tree broadcast or other broadcast that does not rely on pipelining of communication and computations.

#### 10.5. Pipelining multiplications

Another interesting observation is that, if a number of matrix–matrix multiplications need to be performed, the communication and computation can be pipelined *between individual multiplications*.

### 11. SUMMA, ScaLAPACK, AND DISTRIBUTED BLAS

We believe that the SUMMA approach is particularly appropriate for implementation of distributed BLAS implementations of the matrix–matrix multiplication. We summarize those in this Section.

It is very interesting to note that we started pursuing the presented algorithm by making the following simple observation. The blocked right-looking  $LU$  factorization, as im-



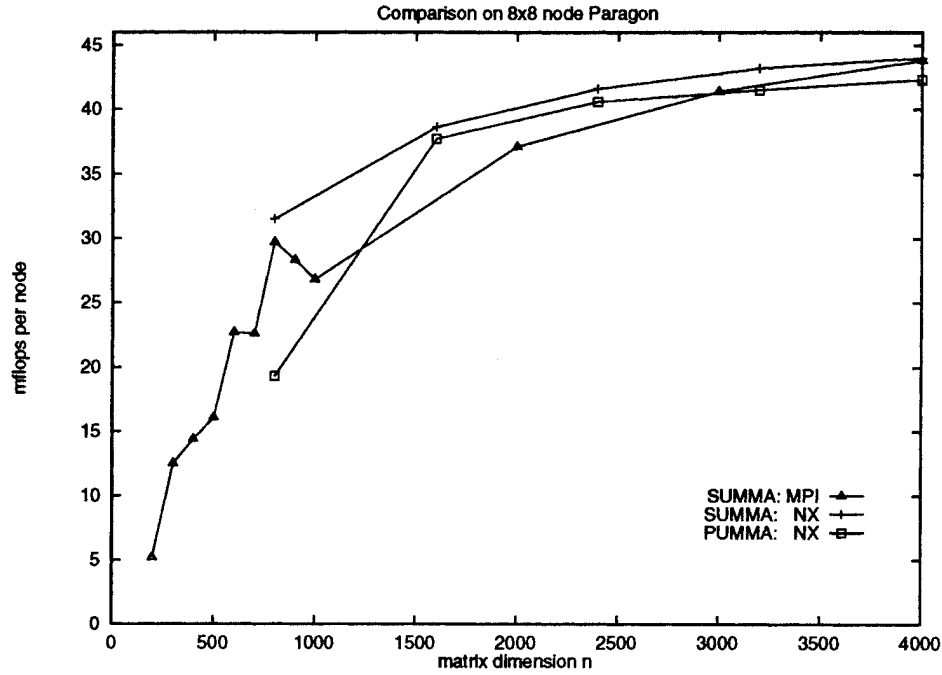


Figure 13. Performance of SUMMA vs. PUMMA for  $C = AB^T$  on 64 nodes

plemented in LAPACK[18,19], is much like a matrix–matrix multiplication,  $C = AB$ , implemented as a series of rank  $nb$  updates, except that they require pivoting, matrices  $A$ ,  $B$  and  $C$  are all the same matrix, and the updates progressively affect less of the matrix being updated. Similarly, the approach used to derive the algorithm for  $C = AB^T$  was actually inspired by the implementation of a left-looking Cholesky factorization.

We hence suspect that a ScaLAPACK implementation based on a distributed BLAS matrix–matrix multiplication would naturally benefit from SUMMA.

## 12. CONCLUSION

The presented algorithms for matrix–matrix multiplication are considerably simpler than those previously presented, all of which have been based on generalizations of the broadcast-multiply-roll algorithm. Nonetheless, their performance is competitive or better, and they are considerably more flexible. Finally, their memory use for work arrays is much lower than those of the broadcast-multiply-roll algorithm. As a result, we believe the SUMMA approach to be the natural choice for a general-purpose implementation.

We should note that on some systems PUMMA may very well outperform SUMMA under some circumstances. In particular, SUMMA is slightly more sensitive to communication overhead. However, it is competitive, or faster, and, given its simplicity and flexibility, warrants consideration. Moreover, the implementations by Huss-Lederman *et al.*[7,8] are competitive with PUMMA, and would thus compare similarly with SUMMA. Also, our




Figure 14. Performance of SUMMA (NX version) for  $C = AB^T$  as a function of the number of nodes, when memory use per node is held constant. The curves for  $100 \times 100$ ,  $200 \times 200$ , etc., indicate the use of local memory equivalent to that necessary to store a matrix of the indicated size, e.g., on 64 and 128 nodes,  $100 \times 100$  is equivalent to a  $800 \times 800$  and  $1128 \times 1128$  global matrix, respectively

method is presented in a slightly simplified setting and thus the performance may be slightly better than it would be if we implemented exactly for the cases for which PUMMA and the algorithm by Huss-Lederman *et al.* were designed.

### 12.1. Further information

We maintain a web page on work related to this paper, which can be accessed at the address

<http://www.cs.utexas.edu/users/plapack>

### ACKNOWLEDGEMENTS

This research was performed in part using the Intel Paragon System and the Intel Touchstone Delta System operated by the California Institute of Technology on behalf of the Concurrent Supercomputing Consortium. Access to this facility was provided by Intel Scalable Systems Division and the California Institute of Technology. This work is partially supported by the NASA High Performance Computing and Communications Program's Earth and Space Sciences Project under NRA Grant NAG5-2497. Additional support came from the Intel Research Council. Jerrell Watts is being supported by an NSF Graduate Research Fellowship.

## REFERENCES

1. G. H. Golub and C. F. Van Loan, *Matrix Computations*, Johns Hopkins University Press, 2nd edn., 1989.
2. L. E. Cannon, *A Cellular Computer to Implement the Kalman Filter Algorithm*, Ph.D. thesis, Montana State University, 1969.
3. G. C. Fox, M. A. Johnson, G. A. Lyzenga, S. W. Otto, J. K. Salmon and D. W. Walker, *Solving Problems on Concurrent Processors*, vol. 1, Prentice Hall, Englewood Cliffs, N.J., 1988.
4. G. Fox, S. Otto and A. Hey, 'Matrix algorithms on a hypercube I: matrix multiplication', *Parallel Comput.*, **3**, 17–31 (1987).
5. C. Lin and L. Snyder, 'A matrix product algorithm and its comparative performance on hypercubes', in Q. Stout and M. Wolfe (Eds.) *Proceedings of Scalable High Performance Computing Conference*, IEEE Press, Los Alamitos, CA, 1992, pp. 190–193.
6. J. Choi, J. J. Dongarra and D. W. Walker, 'PUMMA: Parallel universal matrix multiplication algorithms on distributed memory concurrent computers', *Concurrency, Pract. Exp.*, **6**, (7), 543–570 (1994).
7. S. Huss-Lederman, E. Jacobson and A. Tsao, 'Comparison of scalable parallel matrix multiplication libraries', in *Proceedings of the Scalable Parallel Libraries Conference*, Starksville, MS, Oct. 1993.
8. S. Huss-Lederman, E. Jacobson, A. Tsao and G. Zhang, 'Matrix multiplication on the Intel Touchstone DELTA', *Concurrency, Pract. Exp.*, **6**, (7), 571–594 (1994).
9. W. Gropp, E. Lusk and A. Skjellum, *Using MPI: Portable Programming with the Message-Passing Interface*, The MIT Press, 1994.
10. R. C. Agarwal, F. Gustavson and M. Zubair, 'A high-performance matrix multiplication algorithm on a distributed memory parallel computer using overlapped communication', *IBM J. Res. Dev.*, **38**, (6), (1994).
11. K. Klimkowski and R. van de Geijn, 'Anatomy of a parallel out-of-core dense linear solver', Vol III - Algorithms and applications, *Proceedings of the 1995 International Conference on Parallel Processing*, pp. 29–33.
12. J. J. Dongarra, J. Du Croz, S. Hammarling and I. Duff, 'A set of level 3 basic linear algebra subprograms', *TOMS*, **16**, (1), 1–16 (1990).
13. M. Barnett, S. Gupta, D. Payne, L. Shuler, R. van de Geijn and J. Watts, 'Interprocessor collective communication library (InterCom)', *Scalable High Performance Computing Conference 1994*.
14. M. Barnett, D. G. Payne, R. van de Geijn and J. Watts, 'Broadcasting on meshes with worm-hole routing', *J. Parallel Distrib. Comput.*, **35**, (2), 111–122 (1996).
15. J. Watts and R. van de Geijn, 'A pipelined broadcast for multidimensional meshes', *Parallel Process. Lett.*, **5**, (2), 281–292 (1995).
16. A. Chitchekanova, J. Gunnels, G. Morrow, J. Overfelt and R. van de Geijn, 'Parallel implementation of BLAS: General techniques for level 3 BLAS', *TR-95-40*, Department of Computer Sciences, University of Texas, October 1995.
17. J. J. Dongarra, R. A. van de Geijn and D. W. Walker, 'Scalability issues affecting the design of a dense linear algebra library', *J. Parallel Distrib. Comput.*, **22**, (3), 523–537 (1994).
18. E. Anderson, Z. Bai, C. Bischof, J. Demmel, J. Dongarra, J. DuCroz, A. Greenbaum, S. Hammarling, A. McKenney and D. Sorensen, 'Lapack: A portable linear algebra library for high performance computers', *Proceedings of Supercomputing '90*, IEEE Press, 1990, pp. 1–10.
19. E. Anderson, Z. Bai, J. Demmel, J. Dongarra, J. DuCroz, A. Greenbaum, S. Hammarling, A. McKenney, S. Ostrouchov and D. Sorensen, *LAPACK Users' Guide*, SIAM, Philadelphia, 1992.
20. J. Choi, J. J. Dongarra, R. Pozo and D. W. Walker, 'Scalapack: A scalable linear algebra library for distributed memory concurrent computers', *Proceedings of the Fourth Symposium on the Frontiers of Massively Parallel Computation*, IEEE Comput. Soc. Press, 1992, pp. 120–127.
21. J. Choi, J. J. Dongarra and D. W. Walker, 'Level 3 BLAS for distributed memory concurrent computers', *CNRS-NSF Workshop on Environments and Tools for Parallel Scientific Computing*, Saint Hilaire du Touvet, France, Sept. 7-8, 1992, Elsevier Science Publishers, 1992.
22. J. J. Dongarra, I. S. Duff, D. C. Sorensen and H. A. van der Vorst, *Solving Linear Systems on Vector and Shared Memory Computers*, SIAM, Philadelphia, 1991.

- 
23. C.-T. Ho and S. L. Johnsson, 'Distributed routing algorithms for broadcasting and personalized communication in hypercubes', in *Proceedings of the 1986 International Conference on Parallel Processing*, IEEE, 1986, pp. 640–648.